

Building a **Minimal OS** using Lottery Scheduling

Compact OS for preemptive multitasking

Design and implement a small OS emphasizing interrupts, scheduler, and randomness.

Build a Minimal OS -- LAN-iX

Hands-on kernel basics: scheduling, interrupts, system calls



Modern kernels are complex

1

Paging, VFS, SMP and many subsystems obscure fundamentals

2

Project focus areas

Scheduling, interrupts, randomness, timer preemption

Educational approach

3

Build a minimal but functional OS from scratch to observe core mechanics

Key learning outcomes

4

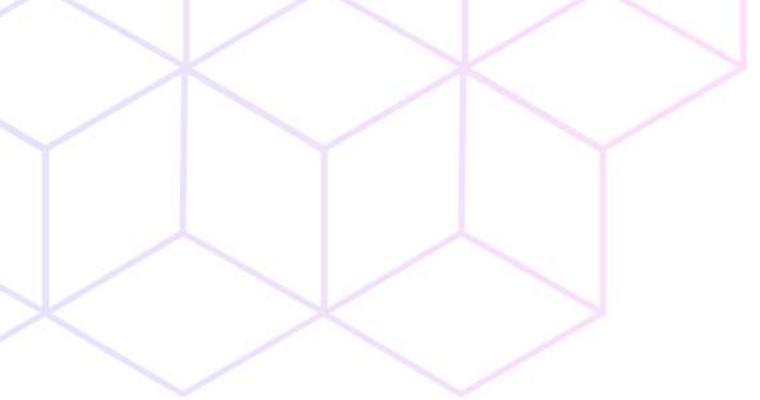
Trap and interrupt flow, protected-mode constraints, interrupt latency

5

Scope trade-offs
Omit advanced features like paging to simplify learning and debugging



Project Scope — Included & Excluded



1 Included features

- Used GRUB bootloader for multiboot support, easier implementation.
- Basic memory model without paging (flat physical memory model for simplicity)
- Simple PCB and context switching to manage task state transitions reliably



2 Included runtime and services

- Lottery Scheduling as the CPU scheduler to explore probabilistic fairness
- Timer interrupts with PIC remapping to enable preemption and proper IRQ routing
- PRNG using CPU timing and interrupt entropy for lightweight randomness support

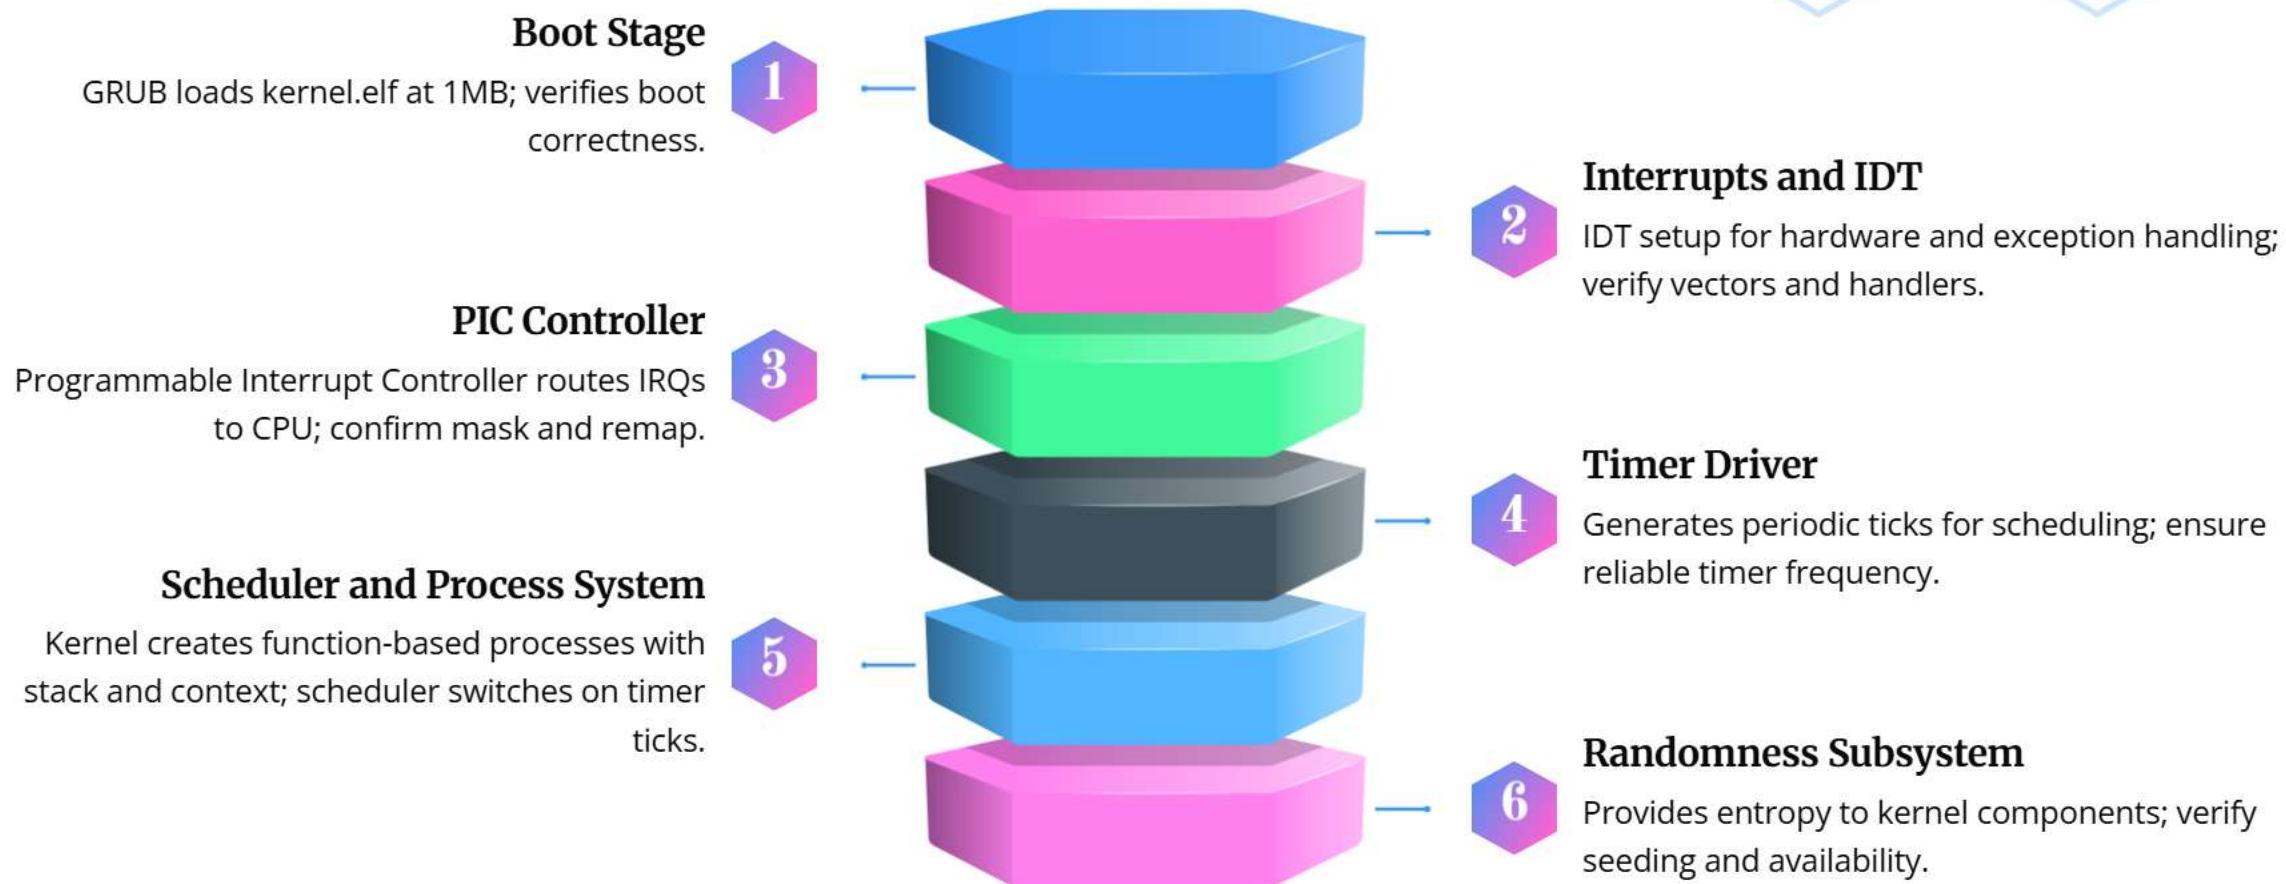


3 Intentionally excluded features

- Paging and virtual memory to avoid the complexity of memory virtualization
- User mode support to keep the kernel/privilege model minimal for the course
- Filesystem implementation to narrow scope away from persistent storage challenges

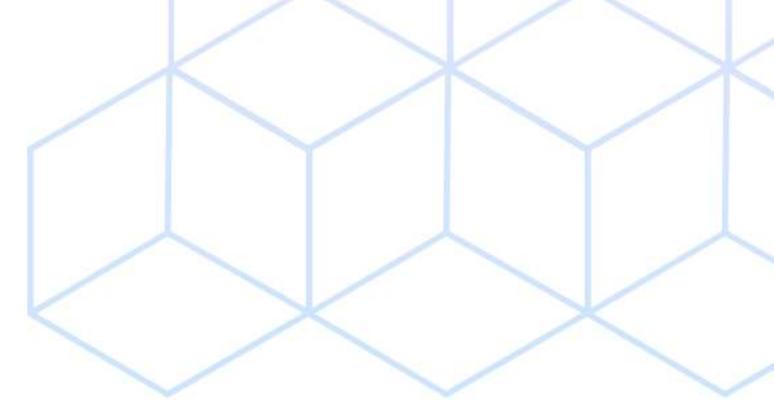
System Architecture Overview

Bootloader to scheduler: control and verification checkpoints



Boot and Entry Code: deterministic early startup

Tiny boot.S sets video memory, switches to protected mode, and jumps to kernel_main



GRUB loads kernel.elf

Verify multiboot header presence, alignment, and entry address at 1MB for GRUB compatibility



boot.S initializes video memory

Confirm video memory writes for early debug output are reliable for diagnosing boot issues

Switch to protected mode

Ensure mode transition sequence is correct and CPU is in protected mode before continuing

Jump to kernel_main

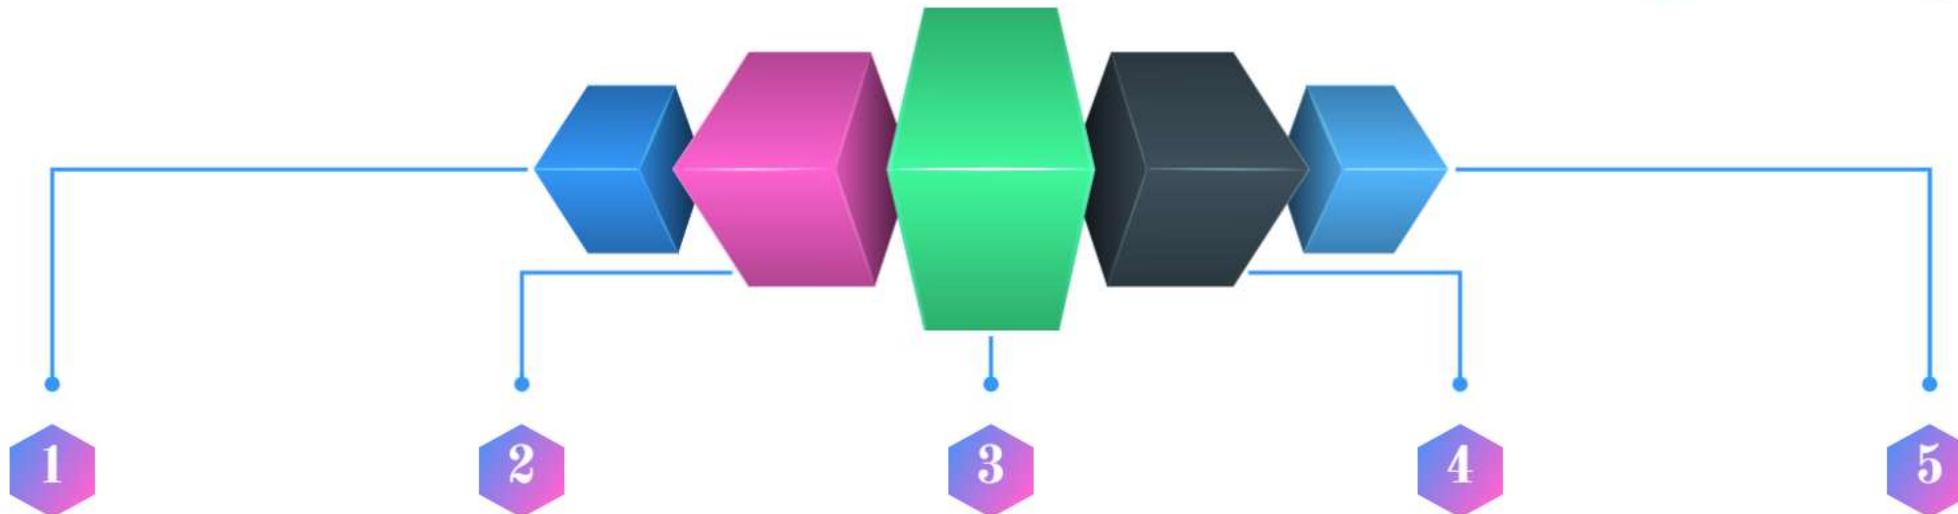
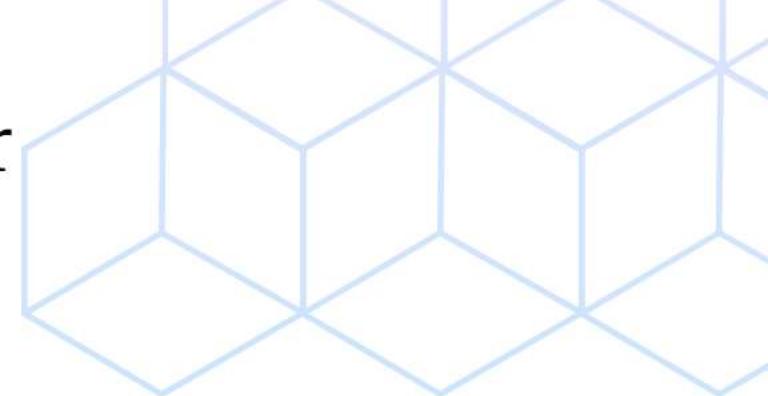
Validate the entry address and that registers/stack meet kernel_main ABI expectations

Initialize IDT

Load IDT early to allow later interrupt based scheduling; verify IDT base and limit

Interrupts and the IDT: Timer, ISR, and Scheduler

How IRQ0 maps to vector 32 and drives preemptive scheduling



1 IDT Setup

IDT configured in idt.c;
load_idt() installs the table
used by all interrupts.

2 IRQ0 to Vector 32

Timer interrupt (IRQ0) is
mapped to vector 32 and
invokes the timer ISR.

3 ISR Entry and Exit

Validate handler entry and
exit sequences and correct
IRQ vector usage.

4 Interrupts Handle

interrupts_handle mixes
timing entropy and forwards
events to the scheduler.

5 Scheduler Interaction

Interrupts are fully functional
and essential for preemptive
scheduling; state changes
must be atomic relative to
scheduler decisions.



Timer Subsystem: 50 Hz PIT Tick

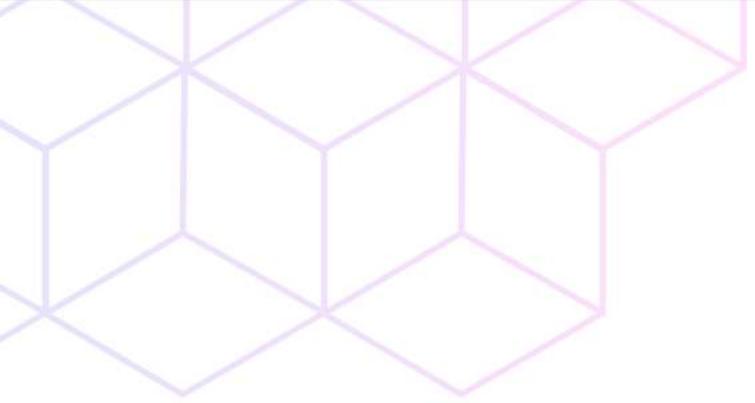
PIT tick flow, scheduler impact, and ISR guidance

- ➊ PIT generates tick
 - ➎ Programmable Interval Timer set to 50 Hz issues a hardware tick
 - ➋ ISR entry
 - ➎ Interrupt service routine runs on tick; keep work minimal to preserve latency
 - ➌ Mix TSC entropy
 - ➎ Add TSC-based entropy into PRNG pool during ISR
 - ➍ Call lottery_schedule_tick
 - ➎ Notify scheduler to consider preemption and update runtime accounting
 - ➎ Send EOI to PIC
 - ➎ Acknowledge interrupt to the PIC so new interrupts can be accepted
 - ➏ Scheduler effect
- ➐ Tick makes OS preemptive; frequency affects responsiveness and overhead
 - ➑ Practical guidance
 - ➎ Balance tick rate for responsiveness versus CPU cost; document fairness and entropy impact

Process Model: fixed-array PCBs

MAX_PROCESSES = 10; each process is a function with its own stack

PID Unique process identifier	State: UNUSED to RUNNABLE to RUNNING Scheduler and interrupts drive transitions	Tickets Used by lottery or weighted scheduler	CPU context (ESP saved/restored) Registers saved on context switch	Stack pointer and stack mapping Each PCB points to its own stack of 1024 words
Fixed array allocation Array size MAX_PROCESSES = 10 simplifies allocation	Process = function with independent stack Function executes using its 1024-word stack	Testing and correctness notes Consider stack sizing, overflow risks, and interrupt-driven transitions	Diagram: PCB maps to stack Visual: PCB entry pointing to its 1024-word stack	



Context Switching: saving CPU state for multitasking

Preserve registers, stack, and flags to resume another process

- ➊ Save registers with pusha
 - ➎ Push all general purpose registers to stack to preserve caller and callee state
- ➋ Save old ESP
 - ➎ Store current stack pointer so the current process can resume with correct stack
- ➌ Load new ESP
 - ➎ Set ESP to the next process stack so its context is active
- ➍ Restore registers with popa
 - ➎ Pop registers from the new stack to restore the next process CPU state
- ➎ Resume execution
 - ➎ Continue execution on the next process stack with preserved EFLAGS and registers
- ➏ Preserve critical state bits
 - ➎ Ensure callee-saved registers and EFLAGS are preserved; validate stack pointers
- ➐ Minimize ISR work
 - ➎ Keep interrupt service routines short to avoid long blocking during context switches

Why use Lottery Scheduling

Probabilistic fairness with minimal code and flexible weights



Lottery Scheduling

- Fairness via ticket-weighted chance
- Processes request more or less CPU by changing tickets
- Small code footprint, no complex priority queues
- Good for experiments and dynamic ticket assignment
- Illustrates probabilistic scheduling concepts for labs



Round Robin and Priority Queues

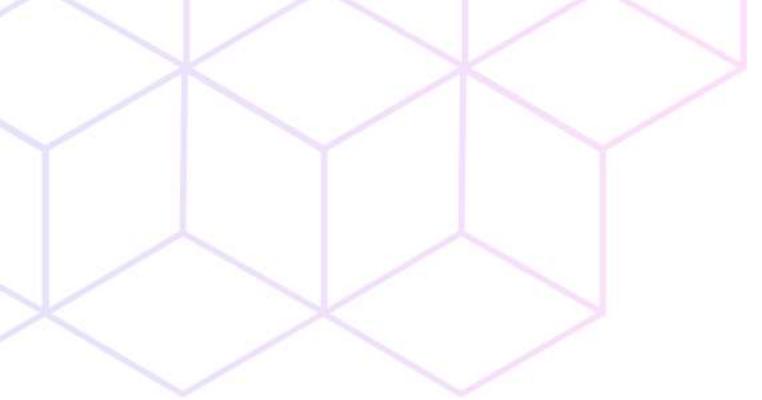
- Round robin: simple and predictable time slices
- Priority queues: deterministic prioritization by numeric priority
- Round robin can be unfair for differing needs
- Priority queues require complex data structures and tuning
- Less flexible for runtime, fine-grained fairness experiments



Deterministic Lottery Scheduling with Atomic Ticket Updates

Fair, preemptive selection using a global ticket pool

- ❖ Compute total tickets
 - ❖ Read global `total_tickets`; ensure atomic read while processes change ticket counts
- ❖ Generate winning ticket
 - ❖ `winning = rng_get_range(total_tickets) + 1;`
RNG supplies nondeterministic entropy
- ❖ Traverse circular PCB list
 - ❖ Iterate PCBs summing tickets until cumulative reaches winning ticket
- ❖ Select winner and prepare switch
 - ❖ Mark chosen PCB; ensure ticket totals remain consistent during handoff
- ❖ Context switch
 - ❖ Perform context switch; scheduler runs preemptively on every timer tick
- ❖ Atomic updates and edge cases
 - ❖ Update ticket totals atomically; handle zero-ticket processes with skip or baseline tickets
- ❖ Fairness validation
 - ❖ Use test harnesses to run many trials and validate distribution qualitatively in class



Kernel Randomness: PRNG Inputs and Mixing

Design, initialization, and quality considerations



Primary entropy inputs

rdtsc Time Stamp Counter and interrupt timing jitter



Ongoing entropy injection

Inject fresh entropy on every interrupt to avoid staleness



Mixer: SplitMix64
Use SplitMix64 as a fast, well-studied mixer for internal state



Quality considerations

Combine high-resolution timers with jitter; avoid sole reliance on short-lived sources



Initialization requirements

Ensure strong initial seed and well-initialized mixer before use in scheduling



Adversarial/pathological cases

Limited interrupt activity or synchronized timers can reduce entropy; plan fallback or conservative use



Result

High-quality pseudo-random generator inside the kernel for lottery scheduling

Future Work: Extend OS with incremental features

Prioritize quick wins, plan for larger isolation efforts



System calls for processes

Add simple kernel call interface to expose services to processes



Dynamic ticket assignment via kernel commands

Allow runtime control of scheduling tickets for flexibility



Cooperative yield system call

Enable voluntary context switches to improve responsiveness



Improved debugging and kernel prints

More sophisticated logs to speed diagnostics and testing



Paging for memory isolation

Larger effort to provide isolation and realistic testing



User mode and privilege levels

Major work to enable protected user space and security